

---

# Cluster Management: Unlocking the Secrets of Cluster API and Its Providers



# Intro

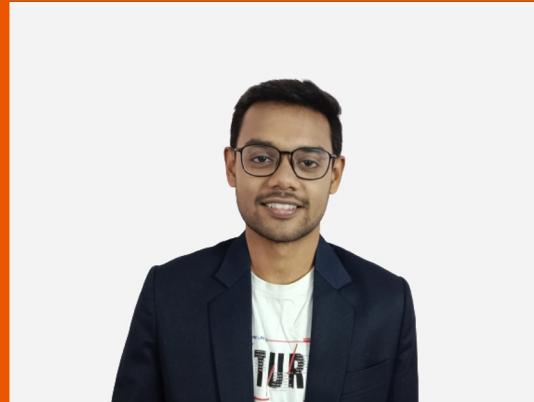


**Subhasmita Swain**

SRE/R&D @CIVO

---

LFX Mentee'22 @CAPG,  
Outreachy'22 @Apache



**Aniruddha Basak**

SWE @Sysself

LFX Mentee'22 @CAPG



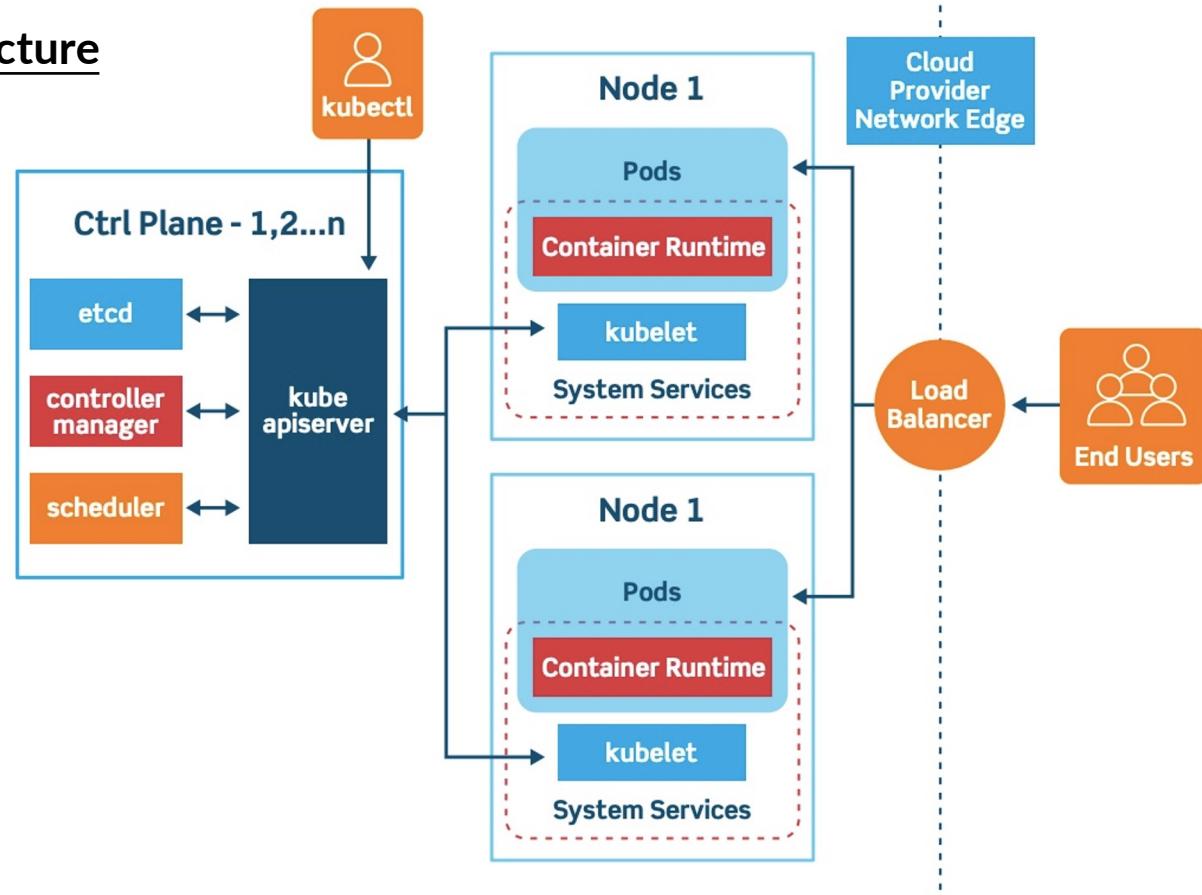
# Agenda

- Cluster Lifecycle
  - Cluster API Architecture
  - Types of CAPI Providers
  - Examples of Infrastructure Providers
-

# All hail KUBERNETES!

Kubernetes provides a framework to run distributed systems resiliently, scaling and managing the lifecycle of containerized applications.

# K8s Cluster Architecture



# Why manage clusters?

1

## Scaling and Load Balancing

- Channel traffic volume
- Maintain uptime & performance
- Distributed load balancing

2

## Automated Rollouts and Rollbacks

- Rollback changes if anything goes wrong
- Progressively rollout changes on configs

3

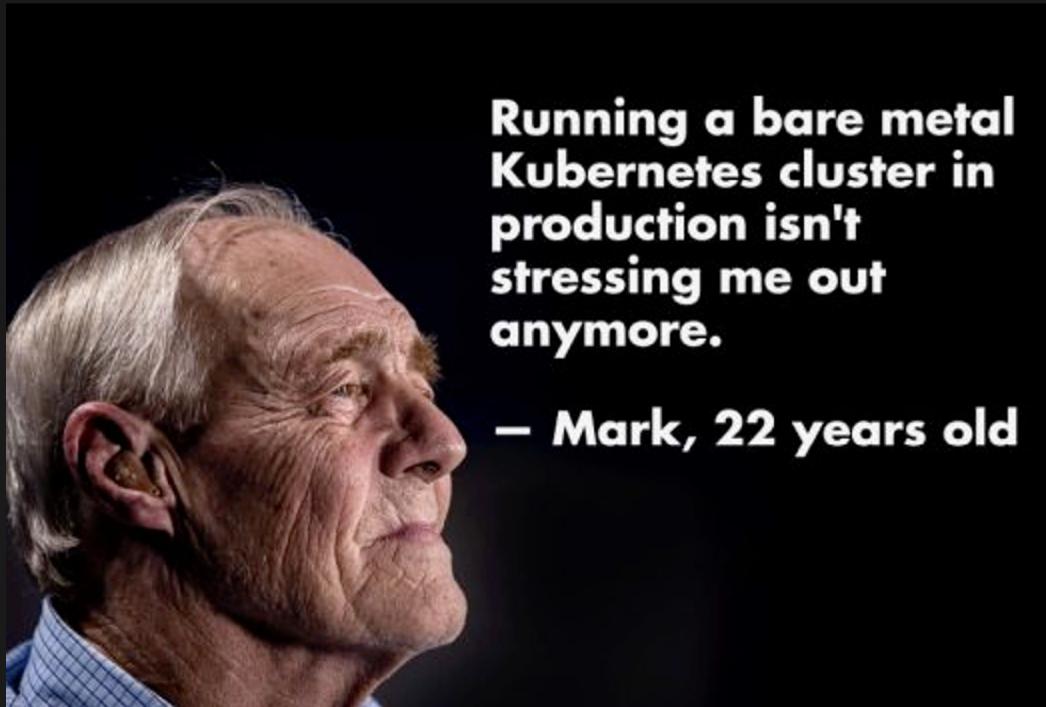
## Self-Healing

- Can detect and replace instances that stop responding
- Based on set params can kill/recreate unresponsive instances

Configuration management & Secrets

4

- Manage sensitive information
- Hierarchical secret masking



**Running a bare metal  
Kubernetes cluster in  
production isn't  
stressing me out  
anymore.**

**– Mark, 22 years old**

# What is CAPI

The Cluster API is a Kubernetes bootstrapping project to bring declarative, Kubernetes-style APIs to multiple cluster creation, configuration, and management.

It provides optional, additive functionality on top of core Kubernetes to manage the lifecycle of a Kubernetes cluster.





---

CAPI, in its simplest form, is a K8s operator that implements **infrastructure-specific** functionality used in conjunction with the main Cluster API when controlling the lifetime of clusters.

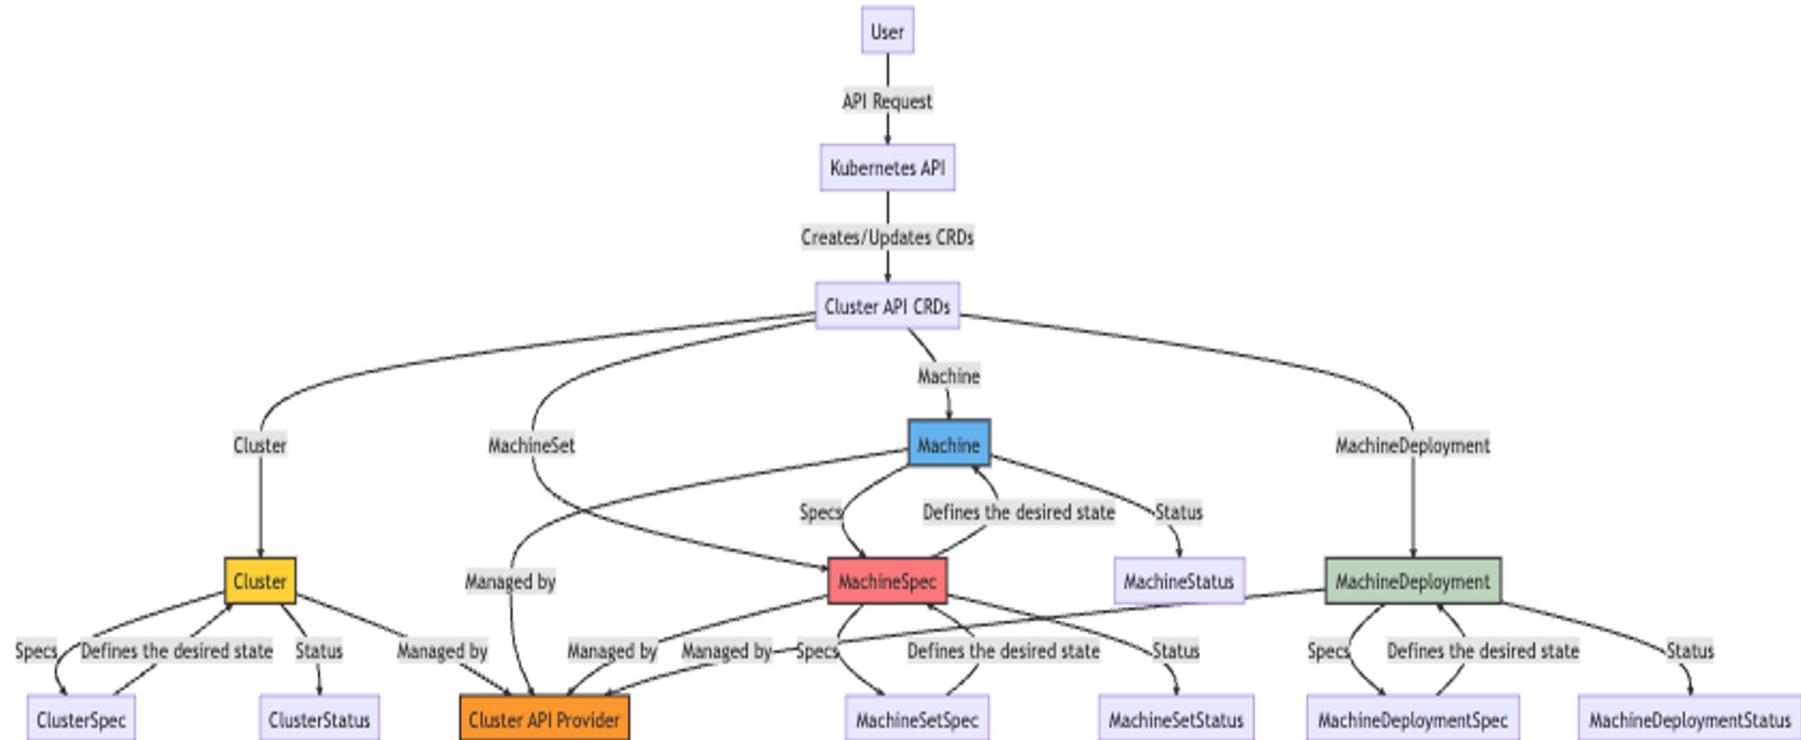
Depending on the kind of provider, the operator implements a contract using its **CRDs (Custom Resources)**.



Today, over **Kubernetes distributions and installers** have been created, each with different default configurations for clusters and supported infrastructure providers.

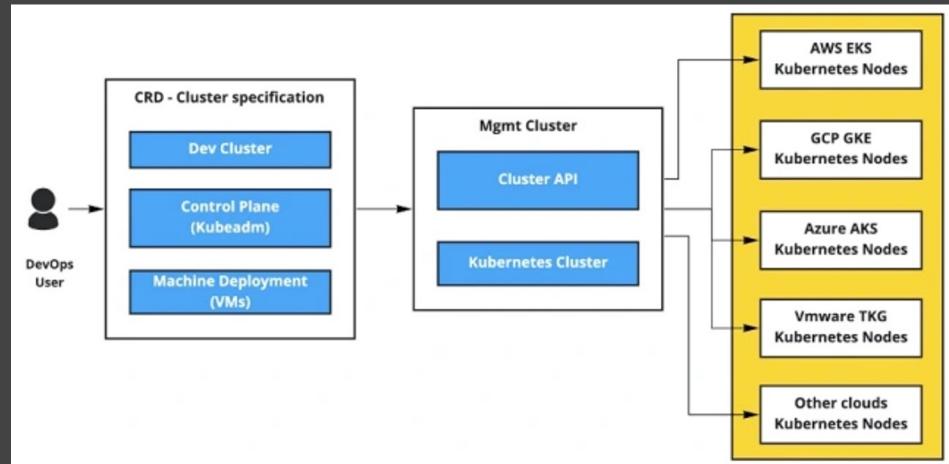


# Understanding the architecture



# CAPI Provider Types

- 
- 1. Bootstrap
  - 2. Control Plane
  - 3. Infrastructure
  - 4. IP Address Management (IPAM)
  - 5. API Adopters and Addons



# Bootstrap

## O1

It is responsible for initializing a machine with the necessary software to join a Kubernetes cluster.

- **Kubernetes Configurations:** generate certificates, setup network configs, install soft.
- **Declarative API:** Defined desired state of cluster
- **Integration with Infrastructure Providers**

Eg., **Kubeadm** sets up CP and worker nodes, a bootstrap provider. Similarly, EKS, Talos, MicroK8s

```
1 apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
2 kind: KubeADMConfigTemplate
3 metadata:
4   name: "test-cluster-md-0"
5 spec:
6   template:
7     spec:
8       joinConfiguration:
9         nodeRegistration:
10        kubeletExtraArgs:
11          cloud-provider: gce
```



# Control Plane

## 02

It is used to control the creation & lifecycle of the kubernetes control plane. It can utilize the resources created by bootstrap and infrastructure providers.

- Kubeadm control plane (KCP) is core, others include
  - MicroK8s, Talos, Nested.
- Managed K8s (i.e., AKS, GKS) implementations
- Can manage multiple CP nodes to ensure high availability of k8s clusters.

```
1 kind: KubeADMControlPlane
2 apiVersion: controlplane.cluster.x-k8s.io/v1beta1
3 metadata:
4   name: "test-control-plane"
5 spec:
6   replicas: 3
7   machineTemplate:
8     infrastructureRef:
9       kind: GCPMachineTemplate
10    apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
11    name: "test-control-plane"
12   kubeADMConfigSpec:
13     initConfiguration:
14       nodeRegistration:
15         kubeletExtraArgs:
16           cloud-provider: gce
17   clusterConfiguration:
18     apiServer:
19       extraArgs:
20         cloud-provider: gce
21     controllerManager:
22       extraArgs:
23         cloud-provider: gce
24     joinConfiguration:
25       nodeRegistration:
26         kubeletExtraArgs:
27           cloud-provider: gce
28   version: "1.26"
```



# Infrastructure

## 03

Responsible for creating, managing and deleting infra resources. Eg., networking, security groups, hosts etc.

- Provider specific implementations.
- Integration with other providers.
- Widely used infra providers include - AWS (CAPA), Azure (CAPZ), GCP (CAPG), etc.

```
● ● ●

1 apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
2 kind: GCPCluster
3 metadata:
4   name: "test-cluster"
5 spec:
6   project: "test-project"
7   region: "us-west1-a"
8   network:
9     name: "test-network"
10 ---
11 kind: GCPMachineTemplate
12 apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
13 metadata:
14   name: "test-cluster-control-plane"
15 spec:
16   template:
17     spec:
18       instanceType: "n2-standard-4"
19       image: "123456789"
20 ---
21 apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
22 kind: GCPMachineTemplate
23 metadata:
24   name: "test-cluster-md-0"
25 spec:
26   template:
27     spec:
28       instanceType: "n2-standard-4"
29       image: "123456789"
```

# IP Address Management (IPAM)

## 04

It is a method of tracking and modifying the information associated with a network's Internet Protocol address space.

In the context of Cluster API (CAPI), an IPAM provider is responsible for managing IP address allocation for clusters, conflict avoidance.

```
1 apiVersion: ipam.example.com/v1
2 kind: IPAMProvider
3 metadata:
4   name: gcp-ipam-provider
5   namespace: kube-system
6 spec:
7   projectID: your-gcp-project-id
8   region: us-central1
9   network: your-gcp-network
10  subnetwork: your-gcp-subnetwork
11  ipRange: 10.0.0.0/16
```



# API Adopters and Addons

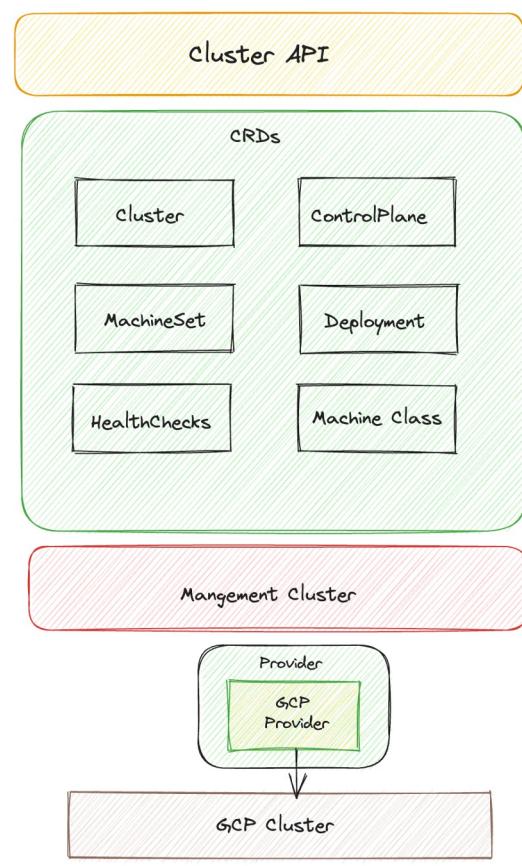
05

API Adopters are implementations managed by 3rd-parties according to the standard CAPI. Eg., Gardener Machine Controller manager, Kubermatic machine controller.

Addons are used to manage the lifecycle of workloads on the cluster after the initial provisioning. Eg., Helm

# Cluster API Provider GCP (CAPG)

Enables to manage both GKE and Compute Engine-based Kubernetes clusters using the same API, providing a consistent experience across different types of GCP-based Kubernetes clusters.

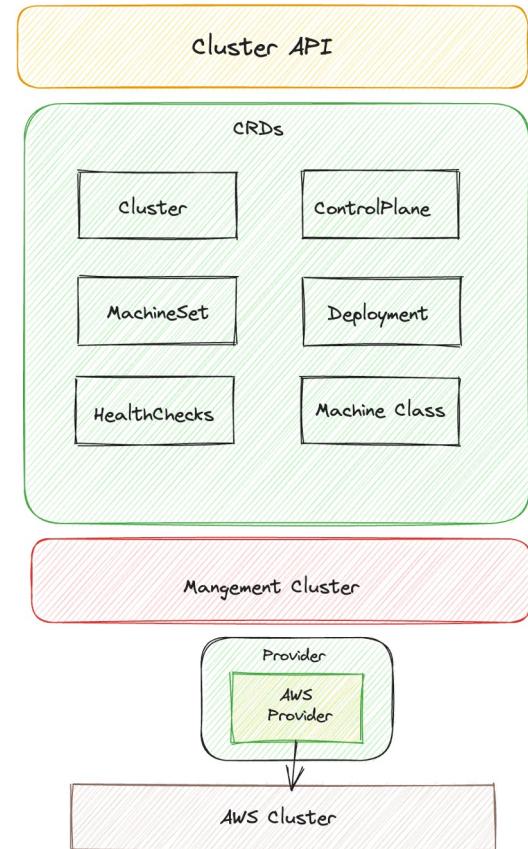


# Cluster API Provider AWS (CAPA)

CAPA offers a Bastion Host feature, enabling secure SSH access to a private network from an external network.

Enabling this feature allows for debugging and maintenance purposes, providing a unique key to the cluster.

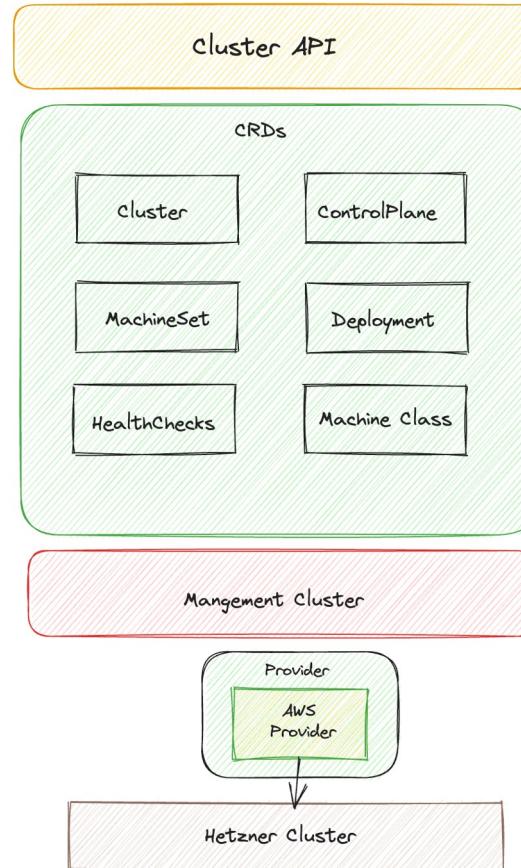
It's like having a special key to your castle!



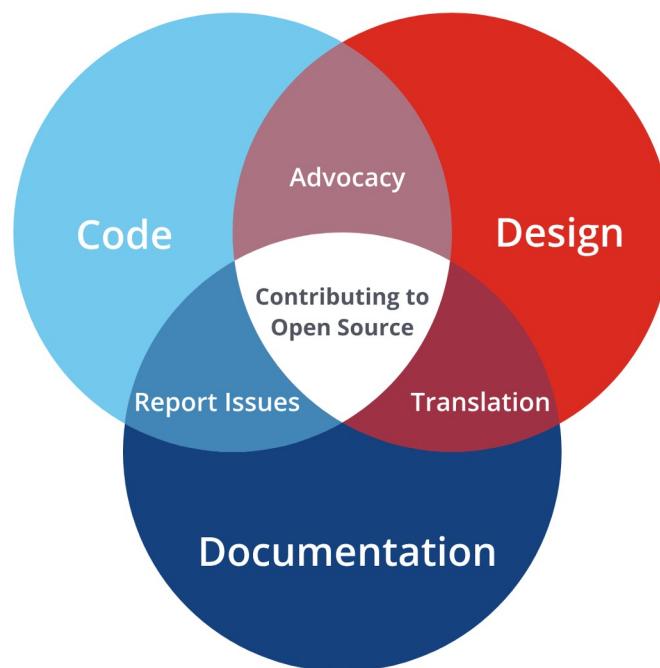
# Cluster API Provider Hetzner (CAPH)

(CAPH) is not only used to manage Kubernetes clusters on Hetzner but it also plays a key role in the testing of Kubernetes itself.

critical component in maintaining the reliability of Kubernetes on Hetzner bare metal infrastructure.



# Contribute Please!



# Resources

Kubernetes Slack



<http://slack.k8s.io/>

Github



<https://github.com/kubernetes-sigs/cluster-api>

CAPI Cookbook



<https://cluster-api.sigs.k8s.io/>

*That's all Folks!*  
*Any Question?*



---

Thank you.

